The background features a subtle, abstract graphic composed of several overlapping circles in shades of blue, green, and yellow, set against a dark blue grid.

Querying in Druid

Jay Wang

14 April 2013

CPS 394 Research Independent Study

Files Involved

- **data.txt**
- **query.body**
- **realtime.spec**
- **start-firehose.sh**
- **run-queries.sh**

data.txt

- Example in our project: appevents.txt
- This is where you specify the structure of your data
- You can use a template described [here](#) or you can just enter in the raw data you want firehose to push
- Sample data stream:

```
{"appid": 1,"event": "clicked ","event_timestamp": 1366079510443,"country": "CA "}
 {"appid": 8,"event": "search ","event_timestamp": 1366079515543,"country": "US"}
 {"appid": 4,"event": "like","event_timestamp": 1366079520643,"country": "US"}
 {"appid": 1,"event": "search ","event_timestamp": 1366079525743,"country": "US"}
```

query.body

- Example in our project: event_counts_query.body
- This is where you define your query
- Here you specify the queryType, dataSource, granularity, dimensions, aggregations, post-aggregations, and intervals
- Specific information to be found [here](#)

Trial I

- Firehose has been set to generate @ 10/sec
- Here is what the data set looks like:

```
 ${name:"sample",type:"combine",endToken:"--",trim:true,separator:""}  
{  
    "appid": ${type:"minmax", min:1, max:10},  
    "event": "${event_name}",  
    "event_timestamp": ${type:"time", format:"millis"},  
    "country": "${country}"  
}  
--  
  
 ${name:"country", type:"weighted",delimiter:",",}  
80:US  
30:CA  
  
 ${name:"event_name",type:"weighted"}  
20:like  
25:search  
18:clicked
```

Trial I cont.

- Here is what the query looks like:

```
{  
  "queryType": "groupBy",  
  "dataSource": "appevents",  
  "granularity": "all",  
  "dimensions": ["appid", "event"],  
  "aggregations": [  
    {"type": "count", "name": "eventcount"},  
    {"type": "doubleSum", "fieldName": "events", "name": "eventssum"}  
],  
  "intervals": ["2012-10-01T00:00/2020-01-01T00"]  
}
```

- Here is what the schema looks like:

```
[  
  {  
    "schema": { "dataSource": "appevents",  
               "aggregators": [ {"type": "count", "name": "events"}  
                ],  

```

Trial I results



Here is a sampling of the results:

```
[ {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "1",  
        "event" : "clicked ",  
        "eventcount" : 4,  
        "eventssum" : 9.0  
    }  
, {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "1",  
        "event" : "like",  
        "eventcount" : 3,  
        "eventssum" : 8.0  
    }  
, {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "1",  
        "event" : "search ",  
        "eventcount" : 3,  
        "eventssum" : 11.0  
    }  
, {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "10",  
        "event" : "clicked ",  
        "eventcount" : 4,  
        "eventssum" : 11.0  
    }  
, {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "10",  
        "event" : "like",  
        "eventcount" : 4,  
        "eventssum" : 9.0  
    }  
, {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "10",  
        "event" : "search ",  
        "eventcount" : 4,  
        "eventssum" : 15.0  
    }  
, {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "2",  
        "event" : "clicked ",  
        "eventcount" : 3,  
        "eventssum" : 15.0  
    }  
, {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "2",  
        "event" : "like",  
        "eventcount" : 3,  
        "eventssum" : 6.0  
    }  
, {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "2",  
        "event" : "search ",  
        "eventcount" : 3,  
        "eventssum" : 9.0  
    }  
, {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "3",  
        "event" : "clicked ",  
        "eventcount" : 2,  
        "eventssum" : 5.0  
    }  
, {
```

Questions

- Q: Why will eventssum \geq eventcount?
A: It lies in how Druid handles aggregations.
- In short, eventssum is the number of events with the specified dimensions
- eventcount is the number of aggregated rows in memory
- If Druid gets duplicate data entities, it will aggregate them into one row and assign a count to that row
 - The count is the number of entities
 - We are using a count aggregator as defined in the schema

Trial II

- Objective: Feed Druid the same data to make the eventcount equal 1.
- New data set:

```
 ${name:"sample",type:"combine",endToken:"----",trim:true,separator:""}  
 {  
     "appid": ${type:"minmax", min:1, max:1},  
     "event": "${event_name}",  
     "event_timestamp": ${type:"time", format:"millis"},  
     "country": "${country}"  
 }  
 --  
  
 ${name:"event_name",type:"weighted"}  
 20:like  
  
 ${name:"country", type:"weighted",delimiter:",",}  
 80:US
```

- As can be seen, every event will have an appid of “1”, event of “like” and country of “US”

Trial II results

- The query and schema will remain the same.
- Results:

```
[ {  
    "version" : "v1",  
    "timestamp" : "-292275055-05-16T16:47:04.192Z",  
    "event" : {  
        "appid" : "1",  
        "event" : "like",  
        "eventcount" : 2,  
        "eventssum" : 230.0  
    }  
} ]
```

- Is this what we would expect?
- Why does the eventcount = 2?

Questions

- ➊ The results make sense.
- ➋ If you go back to the schema for this query, the indexGranularity is set at “minute”
 - ➌ This means at the end of every minute, Druid will perform an aggregation no matter what
 - ➌ Ultimately, this means that the timestamps for these data bits span two minutes
- ➌ Note: Druid aggregates on all of the data’s fields, NOT on dimensions
 - ➌ If country varied from [“US”, “CA”, “EN”], Druid would aggregate separately for {appid:1, event:like, country:US} and {appid:1, event:like, country:CA}, etc.

Stress Tests

- ➊ Trying to write a stress test for Druid
- ➋ “stress” is based on two things:
 - ➌ the amount of data to process over
 - ➌ the complexity of the query
- ➌ Problem:
 - ➌ there is a cap as to how fast firehose can push data to kafka (it’s generally around <1.5K/sec)
 - ➌ i.e. you can’t feasibly push 1M tweets/sec

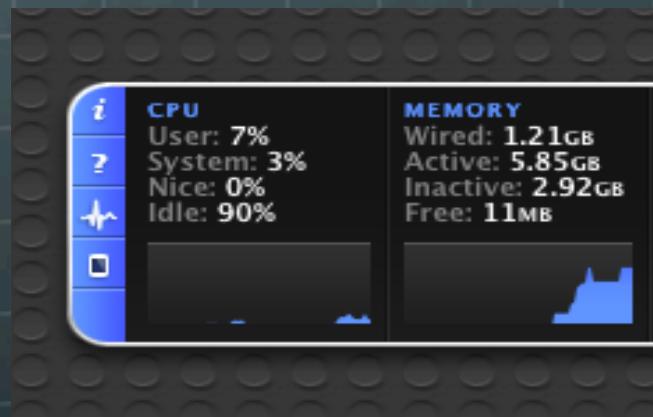
Trial III

- Firehose rate set to 1M/sec (won't actually reach that speed, it peaks out much sooner)
- I will let firehose run for 30 min
- Data file is tweets.txt
- Sample Tweet:

```
{"time_zone":"Alaska","event_timestamp":${type:"time",  
format:"millis"}, "id":306771720193339393, "id_str":"306771720193339393", "text":"Geeee", "source":">\u003ca href=\"http://twitter.  
com/download/iphone\" rel=\"nofollow\"\u003eTwitter for iPhone\u003c/a\u003e", "truncated":false, "in_reply_to_status_id":null, "in_reply_to_status_id_  
str":null, "in_reply_to_user_id":null, "in_reply_to_user_id_str":null, "in_reply_to_screen_name":null, "user": {"id":274990882, "id_str": "274990882", "name":  
"Anne Curtis-Smith", "screen_name": "Pawlavableeee", "location": "Sa ref", "url": "http://johannapawla.tumblr.com/", "description": "Fil-Aussie  
Actress/Host and Pursuer of dreams", "protected": false, "followers_count": 423, "friends_count": 145, "listed_count": 2, "event_timestamp": "Thu Mar 31  
12:26:48 +0000 2011", "favourites_count": 5063, "utc_offset": -32400, "geo_enabled": true, "verified": false, "statuses_count": 23038, "lang": "en", "contributors_  
enabled": false, "is_translator": false, "profile_background_color": "FFFFFF", "profile_background_image_url": "http://si0.twimg.  
com/profile_background_images/779953914/44cd05b6ca02282caecf5656ed80dd16.png", "profile_background_image_url_https": "https://si0.twimg.  
com/profile_background_images/779953914/44cd05b6ca02282caecf5656ed80dd16.png", "profile_background_tile": true, "profile_image_url": "http://si0.twimg.  
.com/profile_images/3300953457/fce4b8629cf7c31aa7547c83723f8097_normal.jpeg", "profile_image_url_https": "https://si0.twimg.  
.com/profile_images/3300953457/fce4b8629cf7c31aa7547c83723f8097_normal.jpeg", "profile_banner_url": "https://si0.twimg.com/profile_banners/2749908  
82/1361611207", "profile_link_color": "#5950FA", "profile_sidebar_border_color": "FFFFFF", "profile_sidebar_fill_color": "FFFFFF", "profile_text_color": "#0000  
00", "profile_use_background_image": true, "default_profile": false, "default_profile_image": false, "following": null, "follow_request_sent": null, "notificatio  
ns": null}, "geo": {"type": "Point", "coordinates": [14.77285192, 120.88028117]}, "coordinates": {"type": "Point", "coordinates": [120.88028117, 14.77285192]}  
}, "place": null, "contributors": null, "retweet_count": 0, "entities": {"hashtags": [], "urls": [], "user_mentions": []}  
}, "favorited": false, "retweeted": false, "filter_level": "medium"}
```

Trial III cont.

- Running the same query, except my dimensions will be [“id”, “id_str”, “text”]
- Schema will remain the same
- I ended up running firehose for 15.4 minutes (my computer began running out of memory)



Stress Success

- Firehose ended up pushing 1.2M tweets

```
generated 1251192 lines in 923895ms (using the System Clock), 1354/s.  
bytes/sec: 3733494, bytes/min: 224009673
```

- Queries ended up taking almost 14 seconds

```
real    0m13.941s  
user    0m0.006s  
sys    0m0.004s  
bin/run_client.sh finished
```

- firehose.log and queries.log file available
- Conclusion: definitely possible to “stress druid”

Issues

- Druid has limited ability to query over nested dimensions
- If still running and the data file has finished processing, firehose will start over (should be easy fix)
- No way to pause firehose when data file is finished and then restart firehose when more data is added to data file

Moving Forward

- For these examples, I have used:
 - queryType: groupBy
 - there also exist: search, timeSeries, timeBoundary
 - aggregations: count and doubleSum
 - there also exist: minMax, longSum
- There are other cool features as well all detailed [here](#)